

BINF5007

Lecture 9: Pandas

@Copyright Samaneh Gholami

Learning Objectives

- 1.Pandas Data Structures:** Understand and differentiate between *Series* and *DataFrame*.
- 2.Pandas Integration:** Utilize pandas with libraries like NumPy and Matplotlib.
- 3.DataFrame Basics:** Create, modify, and explore DataFrames.
- 4.Statistical Analysis:** Compute basic statistics on numerical data.
- 5.Data Import/Export:** Read and write data to formats like Excel and CSV.
- 6.Data Subsetting:** Extract rows, columns, and subsets using indexing and filtering.

The two primary data structures of pandas

- Series (1-dimensional) and DataFrame (2-dimensional)

Dimensions	Name	Description
1	Series	1D labeled homogeneously-typed array
2	DataFrame	General 2D labeled, size-mutable tabular structure with potentially heterogeneously-typed column

- These handle the majority of use cases in finance, statistics, social science, and many areas of engineering
- For R users, the DataFrame provides everything that R's data.frame provides and more
- pandas is built on top of NumPy (see on your own section) and is intended to integrate well within a scientific computing environment with many other 3rd party libraries


Wes McKinney and the pandas Development Team

Why more than one data structure?

- The best way to think about the pandas data structures is as flexible containers for lower dimensional data
- For example, DataFrame is a container for Series, and Series is a container for scalars
- We can insert and remove objects from these containers in a dictionary-like fashion
- pandas has sensible default behaviors for the common API functions which take into account the typical orientation of time series and cross-sectional data sets


pandas Integration

- pandas supports the integration with many file formats or data sources out of the box (csv, excel, sql, json, parquet, etc)
- Importing data from each of these data sources is provided by function with the prefix `read_*`
- Similarly, the `to_*` methods are used to store data



pandas DataFrame

- When working with tabular data like the data stored in spreadsheets or databases, pandas is a good tool
- pandas will help you to explore, clean, and process your data
- **In pandas, a data table is called a DataFrame**
- DataFrames are just another data structure (ds), and we can generate DataFrames using the data structures we learned about last week



DataFrame: 2-dimensional data structures

- It can store data of different types (including characters, integers, floating point values, categorical data and more) in columns
- It is like a spreadsheet, a SQL table or the data frame in R
- This table has 3 columns, each of them with a column label. The column labels are respectively Name, Age, and Sex
- The column Name consists of textual data with each value a string, the column Age are numbers, and the column Sex is textual data

	A	B	C	D
1		Name	Age	Sex
2	0	Braund, Mr. Owen Harris	22	male
3	1	Allen, Mr. William Henry	35	male
4	2	Bonnell, Miss. Elizabeth	58	female
5				

pandas Data Table Representation

- To load the pandas package and start working with it, import the package
- The community agreed alias for pandas is `pd`, so loading pandas as `pd` is assumed standard practice for all of the pandas documentation
- To manually store data in a table, create a **DataFrame**
- When using a Python **dictionary of lists**, the dictionary keys will be used as column headers and the values in each list as columns of the DataFrame

```
<class 'pandas.core.frame.DataFrame'>
```

```
import pandas as pd

df = pd.DataFrame(
    {
        "Name": [
            "Braund, Mr. Owen Harris",
            "Allen, Mr. William Henry",
            "Bonnell, Miss. Elizabeth",
        ],
        "Age": [22, 35, 58],
        "Sex": ["male", "male", "female"],
    }
)
print(df)
```

	Name	Age	Sex
0	Braund, Mr. Owen Harris	22	male
1	Allen, Mr. William Henry	35	male
2	Bonnell, Miss. Elizabeth	58	female

What's this?


Yes, we generate DataFrames using the data structures we learned about last week

Wes McKinney and the pandas Development Team

Thinking About the DataFrame

- With a DataFrame it is more semantically helpful to think of the index (the rows) and the columns rather than axis 0 and axis 1. Iterating through the columns of the DataFrame thus results in more readable code:

```
for col in df.columns:  
    series = df[col]  
    print(series)  
  
0    Braund, Mr. Owen Harris  
1    Allen, Mr. William Henry  
2    Bonnell, Miss. Elizabeth  
Name: Name, dtype: object  
0    22  
1    35  
2    58  
Name: Age, dtype: int64  
0    male  
1    male  
2    female  
Name: Sex, dtype: object
```



Wes McKinney and the pandas Development Team


A Series (Column in a DataFrame)

- Interested in working with the data in the column Age from our previous example?
- When selecting a single column of a pandas DataFrame, the result is a pandas Series
- To select the column, use the column label in between square brackets []

```
print(df["Age"])
0    22
1    35
2    58
Name: Age, dtype: int64
```

The selection of a single column is very similar to selection of dictionary values based on the key

Each column in a DataFrame is a Series
Series



<class 'pandas.core.series.Series'>

Wes McKinney and the pandas Development Team

Working with a DataFrame or Series

- Want to know the maximum Age of the passenger?
- We can do this on the DataFrame by selecting the Age column and applying max():

```
max_age = df["Age"].max()  
print(max_age)  
max_age = ages.max()  
print(max_age)
```

58

58

	A	B	C	D
1		Name	Age	Sex
2	0	Braund, Mr. Owen Harris	22	male
3	1	Allen, Mr. William Henry	35	male
4	2	Bonnell, Miss. Elizabeth	58	female
5				

- As illustrated by the max() method, you can do things with a DataFrame or Series
- pandas provides a lot of functionalities, through methods that you can apply to a DataFrame or Series. As methods are functions, **you need parentheses ()**

Statistics of the Numerical Data

- Getting basic statistics of the numerical data of my data table
- The `describe()` method provides a quick overview of the numerical data in a DataFrame
- As the `Name` and `Sex` columns are textual data, these are by default not considered by the `describe()` method
- Many pandas' operations return a DataFrame or a Series
- The `describe()` method is an example of a pandas operation returning a pandas Series or a pandas DataFrame

```
print(df.describe())
```


	Age
count	3.000000
mean	38.333333
std	18.230012
min	22.000000
25%	28.500000
50%	35.000000
75%	46.500000
max	58.000000

Remember

- Import the package, aka `import pandas as pd`
- A table of data is stored as a pandas DataFrame
- Each column in a DataFrame is a Series
- You can do things by applying a method to a DataFrame or Series

Reading in Tabular Data

How do I read and write tabular data?



- Importing data from each of these data sources is provided by function with the prefix `read_*`. Similarly, the `to_*` methods are used to store data
- pandas provides the `read_csv()` function to read data stored as a csv file into a pandas DataFrame

Getting the Titanic tabular data for this section

- PassengerId: Id of every passenger.
- Survived: This feature have value 0 and 1. 0 for not survived and 1 for survived.
- Pclass: There are 3 classes: Class 1, Class 2 and Class 3.
- Name: Name of passenger.
- Sex: Gender of passenger.
- Age: Age of passenger.
- SibSp: Indication that passenger have siblings and spouse.
- Parch: Whether a passenger is alone or have family.
- Ticket: Ticket number of passenger.
- Fare: Indicating the fare.
- Cabin: The cabin of passenger.
- Embarked: The embarked category.

```
PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C
11,1,3,"Sandstrom, Miss. Marguerite Rut",female,4,1,1,PP 9549,16.7,G6,S
12,1,1,"Bonnell, Miss. Elizabeth",female,58,0,0,113783,26.55,C103,S
```

<https://raw.githubusercontent.com/pandas-dev/pandas/main/doc/data/titanic.csv>

```
wget https://raw.githubusercontent.com/pandas-dev/pandas/main/doc/data/titanic.csv
```

Wes McKinney and the pandas Development Team

Checking The Data Import

- Good idea to check on the data after reading in the data
- When displaying a DataFrame, the first and last 5 rows will be shown by default:

```
import pandas as pd
```

```
titanic = pd.read_csv("data/titanic.csv")
print(titanic)
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Allen, Mr. William Henry	male	35.0	0	0	113803	53.1000	C123	S
4	5	0	3	Montvila, Rev. Juozas	male	27.0	0	0	373450	8.0500	NaN	S
..	Graham, Miss. Margaret Edith	female	19.0	0	0	211536	13.0000	NaN	S
886	887	0	2	Johnston, Miss. Catherine Helen "Cannie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
887	888	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
888	889	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q
889	890	0	3									
890	891	0	3									

- To see the first N rows of a DataFrame, use the `head()` method with the required number of rows (in this case 8) as argument:
`print(titanic.head(8))`
- To see the last N rows: `titanic.tail(10)`

<https://raw.githubusercontent.com/pandas-dev/pandas/main/doc/data/titanic.csv>

Wes McKinney and the pandas Development Team

Check How pandas Interpreted each of the column data types

- Good idea to check the types (later on we'll see how import with typing)

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
print(titanic.dtypes)
```

- For each of the columns, the used data type is enlisted. The data types in this DataFrame are integers (int64), floats (float64) and strings (object)
- When asking for the dtypes, no parenthesis are used, b/c it's an attribute of a DataFrame or Series
- Attributes represent a characteristic of a DataFrame/Series, whereas a method (which requires parenthesis) does something with the DataFrame/Series


PassengerId	int64
Survived	int64
Pclass	int64
Name	object
Sex	object
Age	float64
SibSp	int64
Parch	int64
Ticket	object
Fare	float64
Cabin	object
Embarked	object
dtype:	object

Print Data to Excel in One Line

- Print the Titanic data to an excel spreadsheet

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
titanic.to_excel("titanic.xlsx", sheet_name="passengers", index=False)
```



PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	0	3	Braund, Mr. male	22	1	0	A/5 21171	7.25		S	
2	1	1	Cumings, Mrs. female	38	1	0	PC 17599	71.2833	C85	C	
3	1	3	Heikkinen, female	26	0	0	STON/O2. 3	7.925		S	
4	1	1	Futrelle, Mrs. female	35	1	0	113803	53.1	C123	S	
5	0	3	Allan, Mr. V male	35	0	0	373450	8.05		S	

If you get a: ModuleNotFoundError: No module named 'openpyxl'
Perform a: pipenv install openpyxl # or pip install openpyxl

`read_*` functions are used to read data to pandas, and the `to_*` methods are used to store data. The `to_excel()` method stores the data as an excel file. In the example here, the `sheet_name` is named `passengers` instead of the default `Sheet1`

By setting `index=False` the row index labels are not saved in the spreadsheet

The equivalent `read` function `read_excel()` will reload the data to a DataFrame

Wes McKinney and the pandas Development Team

Technical Information About a DataFrame

- The **method** `info()` provides more technical information about a DataFrame than `dtypes`

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
print(titanic.info())
```

Let's explain the output in more detail: It's a DataFrame

- There are 891 entries, i.e., 891 rows
- Each row has a row label (aka the index) with values ranging from 0 to 890
- The table has 12 columns. **Most** columns have a value for each of the rows (all 891 values are non-null). Some columns do have missing values and less than 891 non-null values
- The columns `Name`, `Sex`, `Cabin` and `Embarked` consists of textual data (strings, aka `object`). The other columns are numerical data with some of them whole numbers (aka `integer`) and others are real numbers (aka `float`)
- The kind of data (characters, integers,...) in the different columns are summarized by listing the `dtypes`
- The approximate amount of RAM used to hold the DataFrame is provided as well

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   PassengerId 891 non-null    int64  
 1   Survived     891 non-null    int64  
 2   Pclass       891 non-null    int64  
 3   Name         891 non-null    object  
 4   Sex          891 non-null    object  
 5   Age          714 non-null    float64 
 6   SibSp        891 non-null    int64  
 7   Parch        891 non-null    int64  
 8   Ticket       891 non-null    object  
 9   Fare         891 non-null    float64 
 10  Cabin        204 non-null    object  
 11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```

Wes McKinney and the pandas Development Team

Remember

- Import the package, aka `import pandas as pd`
- A table of data is stored as a pandas DataFrame
- Each column in a DataFrame is a Series
- You can do things by applying a method to a DataFrame or Series
- Getting data into pandas from many different file formats or data sources is supported by `read_*` functions
- Exporting data out of pandas is provided by different `to_*` methods
- The `head/tail/info` methods and the `dtypes` attribute are convenient for a first check

Getting a Subset from a DataFrame


How do I select specific columns from a DataFrame?

- Selecting specific columns from a DataFrame:

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
ages = titanic["Age"]
print(ages)
print(ages.shape)
```

```
0      22.0
1      38.0
2      26.0
3      35.0
4      35.0
...
886    27.0
887    19.0
888    NaN
889    26.0
890    32.0
Name: Age, Length: 891, dtype: float64
(891,) ←
```



To select a single column, use square brackets [] with the column name of the column of interest

Each column in a DataFrame is a Series. As a single column is selected, the returned object is a pandas Series. **How could you verify?**

DataFrame.shape is an attribute of a pandas Series and DataFrame containing the number of rows and columns: (nrows, ncolumns). A pandas Series is 1-dimensional and only the number of rows is returned

How do I select specific columns from a DataFrame?

- How do I select **multiple** columns from a DataFrame?

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
age_sex = titanic[["Age", "Sex"]]
print(age_sex)
print(age_sex.shape)
```

	Age	Sex
0	22.0	male
1	38.0	female
2	26.0	female
3	35.0	female
4	35.0	male
..
886	27.0	male
887	19.0	female
888	NaN	female
889	26.0	male
890	32.0	male
[891 rows x 2 columns]		
(891, 2)		



To select a multiple columns, use square brackets [] with the column names of interest

When multiple column are selected, the returned object is a pandas DataFrame. **How could you verify?**

`DataFrame.shape` is an attribute of a pandas Series and DataFrame containing the number of rows and columns: (nrows, ncolumns). A pandas DataFrame is 2-dimensional, so the the number of rows and columns is returned


Chesley Leslin

How do I filter specific rows from a DataFrame?

- Interested in the passengers older than 35 years:

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
above_35 = titanic[titanic["Age"] > 35]
print(above_35.head())
```



	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... McCarthy, Mr. Timothy J	female male	38.0 54.0	1 0	0 0	PC 17599	71.2833 51.8625	C85 E46	C S
6	7	0	1	Bonnell, Miss. Elizabeth	female	58.0	0	0	17463	26.5500	C103	S
11	12	1	1	Andersson, Mr. Anders Johan	male	39.0	1	5	113783	31.2750	NaN	S
13	14	0	3	Hewlett, Mrs. (Mary D Kingcome)	female	55.0	0	0	347082	16.0000	NaN	S
15	16	1	2						248706			

To select rows based on a conditional expression, use a condition **inside** the selection brackets **[]**

How do I filter specific rows from a DataFrame?

- Technically speaking, how did the previous expression work?
- The output of the conditional expression (`>`, but also `==`, `!=`, `<`, `<=`, ...) would work) is a pandas Series of boolean values (either `True` or `False`) with the same number of rows as the original DataFrame
- A Series of boolean values can filter the DataFrame by putting it b/t the selection brackets `[]`
- Only rows for which the value is `True` will be selected:
`above_35 = titanic[titanic["Age"] > 35]`
- The original Titanic DataFrame consists of 891 rows. How many rows are in `above_35`?
- What type is `above_35`?

```
>>> titanic["Age"] > 35
Out[14]:
0      False
1      True
2     False
3     False
4     False
...
886    False
887    False
888    False
889    False
890    False
Name: Age, Length: 891, dtype: bool
```

```
>>> above_35.shape[0]
217
>>> len(above_35)
217
>>> type(above_35)
<class 'pandas.core.frame.DataFrame'>
```

Chesley Leslin

DataFrames Can Have NaN Values

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	Nan	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599 STON/O2. 3101282	71.2833 7.9250	C85 NaN	C S
2	3	1	3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0				
3	4	1	1	Allen, Mr. William Henry	male	35.0	1	0	113803	53.1000	C123	S
4	5	0	3		male	35.0	0	0	373450	8.0500	Nan	S
..
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	Nan	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	Nan	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	Nan	Q

How can we work with passenger data for which the age is **known**, i.e. **not NaN** ?

How do I filter specific rows from a DataFrame?

- Working with passenger data for which the age is **known**, i.e. **not NaN**
- The **`notna()`** conditional function returns a True for each row the values are not a Null value, and this can be combined with the selection brackets `[]` to filter the data table

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
age_no_na = titanic[titanic["Age"].notna()]
print(age_no_na.head())
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599 STON/O2. 3101282	71.2833 7.9250	C85 NaN	C S
2	3	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0	113803	53.1000	C123	S
3	4	0	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
4	5	3									

- How can you check how many rows were found? `len(age_no_na)`

How do I select specific rows and columns from a DataFrame

- Interested in the names of the passengers older than 35 years

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
adult_names = titanic.loc[titanic["Age"] > 35, "Name"]
print(adult_names.head())
print(type(adult_names.head()))
```

```
1    Cumings, Mrs. John Bradley (Florence Briggs Th...
6                  McCarthy, Mr. Timothy J
11                 Bonnell, Miss. Elizabeth
13                Andersson, Mr. Anders Johan
15          Hewlett, Mrs. (Mary D Kingcome)
Name: Name, dtype: object
<class 'pandas.core.series.Series'>
```



Using selection brackets [] is **NOT** sufficient anymore for a subset of both **rows** and **columns**

Instead, the `loc/iloc` operators are required in front of the selection brackets []

When using `loc/iloc`, the part **before** the comma is the **rows** you want, and the part **after** the comma is the **columns** you want to select

You can use a list to return multiple columns, but what would it return?

How do I select specific rows and columns from a DataFrame

- For both the part **before** and **after** the comma, you can use a single label, a list of labels, a slice of labels, a conditional expression or a colon. Using a colon specifies you want to select all rows or columns

```
import pandas as pd

titanic = pd.read_csv("data/titanic.csv")
print(titanic.head())
sub_titanic = titanic.iloc[9:25, 2:5] # rows,
columns print(sub_titanic.head())
```



	0	1	2	3	4	5						
	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	Nan	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0	STON/O2. 3101282	7.9250	Nan	S
3	4	1	1	Allen, Mr. William Henry	male	35.0	1	0	113803	53.1000	C123	S
4	5	0	3						373450	8.0500	Nan	S
	Pclass	2:5		Name	Sex							
9:25	2	Nasser, Mrs. Nicholas (Adele Achem)		female								
9	3	Sandstrom, Miss. Marguerite Rut		female								
10	1	Bonnell, Miss. Elizabeth		female								
11	3	Saunderscock, Mr. William Henry		male								
12	3	Andersson, Mr. Anders Johan		male								
13												Chesley Leslin

Interested in rows indexes 9 - 24 and columns 2 – 4

Remember

- When selecting subsets of data, square brackets [] are used
- Inside these brackets, you can use a single column/row label, a list of column/row labels, a slice of labels, a conditional expression or a colon
- Select specific rows and/or columns using `loc` when using the row and column **names**
- Select specific rows and/or columns using `iloc` when using the **positions** in the table
- You can assign new values to a selection based on `loc/iloc`

Your Turn

Find all anonymous passengers in the dataset